

CPU Scheduling

Youngjae Kim (Ph.D.)

Sogang University

Linux 운영체제 및 응용 (GITF315-01)

Scheduling: Introduction

■ Workload assumptions:

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

Scheduling Metrics

- Performance metric: **Turnaround time**

- The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Another metric is **fairness**.

- Performance and fairness are often at odds in scheduling.


First In, First Out (FIFO)

■ First Come, First Served (FCFS)

- Very simple and easy to implement

■ Example:


- A arrived just before B which arrived just before C.
- Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect

- Let's relax assumption 1: Each job no longer runs for the same amount of time.
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = \mathbf{110 \text{ sec}}$$

Passing the Tractor

- **Problem with Previous Scheduler:**

- FIFO: Turnaround time can suffer when short jobs must wait for long jobs

- **New scheduler:**

- SJF (Shortest Job First)
- Choose job with smallest *run_time*


Shortest Job First (SJF)

- Run the shortest job first, then the next shortest, and so on

- Non-preemptive scheduler

- **Example:**


- A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time.
- Example:
 - A arrives at $t=0$ and needs to run for 100 seconds.
 - B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$


Shortest Time-to-Completion First (STCF)

- Add **preemption** to SJF
 - Also known as Preemptive Shortest Job First (PSJF)
- **A new job enters the system:**
 - Determine of the remaining jobs and new job
 - Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

■ Example:

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

New scheduling metric: Response time

- The time from when the job arrives to the first time it is scheduled.

$$T_{response} = T_{firstrun} - T_{arrival}$$

- STCF and related disciplines are not particularly good for response time.

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling


■ Time slicing Scheduling

- Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a scheduling quantum.
- It repeatedly does so until the jobs are finished.
- The length of a time slice must be *a multiple of* the timer-interrupt period.

**RR is fair, but performs poorly on metrics
such as turnaround time**


RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



SJF (Bad for Response Time)

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$



RR with a time-slice of 1sec (Good for Response Time)

$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1sec$$

The length of the time slice is critical.

■ The shorter time slice

- Better response time
- The cost of context switching will dominate overall performance.

■ The longer time slice

- Amortize the cost of switching
- Worse response time

Deciding on the length of the time slice presents
a **trade-off** to a system designer


Incorporating I/O

- **Let's relax assumption 3: All programs perform I/O**


- **Example:**

- A and B need 50ms of CPU time each.
- A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
- B simply uses the CPU for 50ms and performs no I/O
- The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Poor Use of Resources



**Maximize the
CPU utilization**

Overlap Allows Better Use of Resources

Incorporating I/O (Cont.)

- **When a job initiates an I/O request.**
 - The job is blocked waiting for I/O completion.
 - The scheduler should schedule another job on the CPU.

- **When the I/O completes**
 - An interrupt is raised.
 - The OS moves the process from blocked back to the ready state.

The Multi-Level Feedback Queue

Multi-Level Feedback Queue (MLFQ)

- A Scheduler that learns from the past to predict the future.
- Objective:
 - Optimize **turnaround time** → Run shorter jobs first
 - Minimize **response time** without *a priori knowledge of job length.*

MLFQ: Basic Rules

- **MLFQ has a number of distinct queues.**
 - Each queue is assigned a different priority level.
- **A job that is ready to run is on a single queue.**
 - A job on a **higher queue** is chosen to run.
 - Use round-robin scheduling among jobs in the same queue


Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: Basic Rules (Cont.)

- MLFQ varies the priority of a job based on **its observed behavior**.
- **Example:**
 - A job repeatedly relinquishes the CPU while waiting IOs → Keep its priority high
 - A job uses the CPU intensively for long periods of time → Reduce its priority.

MLFQ Example



MLFQ: How to Change Priority


■ MLFQ priority adjustment algorithm:

- **Rule 3:** When a job enters the system, it is placed at the highest priority
- **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
- **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job


- A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job

■ Assumption:


- **Job A:** A long-running CPU-intensive job
- **Job B:** A short-running interactive job (20ms runtime)
- A has been running for some time, and then B arrives at time $T=100$.



Example 3: What About I/O?

■ Assumption:

- **Job A:** A long-running CPU-intensive job
- **Job B:** An interactive job that need the CPU only for 1ms before performing an I/O



A Mixed I/O-intensive and CPU-intensive Workload (msec)

The MLFQ approach keeps an interactive job at the highest priority

Problems with the Basic MLFQ

■ Starvation

- If there are “too many” interactive jobs in the system.
- Lon-running jobs will never receive any CPU time.

■ Game the scheduler

- After running 99% of a time slice, issue an I/O operation.
- The job gains a higher percentage of CPU time.

■ A program may change its behavior over time.


- CPU bound process → I/O bound process

The Priority Boost

- **Rule 5: After some time period S , move all the jobs in the system to the topmost queue.**

- Example:

- A long-running job(A) with two short-running interactive job(B, C)




Without(Left) and With(Right) Priority Boost

A:  B:  C: 

Better Accounting

- How to prevent gaming of our scheduler?
- Solution:
 - **Rule 4** (Rewrite Rules 4a and 4b): Once a job **uses up its time allotment** at a given level (regardless of how many times it has given up the CPU), **its priority is reduced**(i.e., it moves down on queue).




Without(Left) and With(Right) Gaming Tolerance

Tuning MLFQ And Other Issues

Lower Priority, Longer Quanta

- **The high-priority queues → Short time slices**
 - E.g., 10 or fewer milliseconds
- **The Low-priority queue → Longer time slices**
 - E.g., 100 milliseconds



Example) 10ms for the highest queue, 20ms for the middle, 40ms for the lowest

The Solaris MLFQ implementation

- **For the Time-Sharing scheduling class (TS)**
 - 60 Queues
 - Slowly increasing time-slice length
 - The highest priority: 20msec
 - The lowest priority: A few hundred milliseconds
 - Priorities boosted around every 1 second or so.

MLFQ: Summary

■ The refined set of MLFQ rules:

- **Rule 1:** If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
- **Rule 2:** If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.
- **Rule 3:** When a job enters the system, it is placed at the highest priority.
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced(i.e., it moves down on queue).
- **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.

■ Beauty of MLFQ

- It does not require prior knowledge on the CPU usage of a process.

Multiprocessor Scheduling


Multiprocessor Scheduling

- The rise of the **multicore processor** is the source of multiprocessor-scheduling proliferation.
 - **Multicore:** Multiple CPU cores are packed onto a single chip.
- Adding more CPUs does not make that single application run faster. → You'll have to rewrite application to run in parallel, using threads.

How to schedule jobs on **Multiple CPUs?**

Example: Multicore CPU Cache Hierarchy

Processor package



L1 i-cache and d-cache:
32 KB, 8-way,
Access: 4 cycles


L2 unified cache:
256 KB, 8-way,
Access: 10 cycles

L3 unified cache:
8 MB, 16-way,
Access: 40-75 cycles

Block size: 64 bytes for all
caches.

Example: Multicore CPU Cache Hierarchy

Processor package



L1 i-cache and d-cache:
32 KB, 8-way,
Access: 4 cycles


L2 unified cache:
256 KB, 8-way,
Access: 10 cycles

L3 unified cache:
8 MB, 16-way,
Access: 40-75 cycles

Block size: 64 bytes for all
caches.

Example: Multicore CPU Cache Hierarchy

Processor package



L1 i-cache and d-cache:
32 KB, 8-way,
Access: 4 cycles


L2 unified cache:
256 KB, 8-way,
Access: 10 cycles

L3 unified cache:
8 MB, 16-way,
Access: 40-75 cycles


Block size: 64 bytes for all
caches.

Single Queue Multiprocessor Scheduling (SQMS)

- Put all jobs that need to be scheduled into **a single queue**.
 - Each CPU simply picks the next job from the globally shared queue.




Possible job scheduler across CPUs:



Single Queue Multiprocessor Scheduling (SQMS)


- Put all jobs that need to be scheduled into **a single queue**.
 - Each CPU simply picks the next job from the globally shared queue.



Synchronization Overhead:


- Multiple CPUs attempting to access and modify the same queue concurrently create **race conditions**.
- To ensure data integrity, **locks** (mutexes) must be used.

Possible job scheduler across CPUs:



Single Queue Multiprocessor Scheduling (SQMS)


- Put all jobs that need to be scheduled into a **single queue**.
 - Each CPU simply picks the next job from the globally shared queue.



Synchronization Overhead:

- Multiple CPUs attempting to access and modify the same queue concurrently create **race conditions**.
- To ensure data integrity, **locks** (mutexes) must be used.

Possible job scheduler across CPUs:






Frequent Process Migration:

- A process can run on CPU0, but might be executed on CPU1 for its next turn.

Cache Misses and Perf. Degradation:

- When the process moves to a new CPU (CPU1), its cached data is not present in CPU1's cache.

Scheduling Example with Cache Affinity




- Preserving affinity for most
 - Jobs A through D are not moved across processors.
 - Only job e Migrating from CPU to CPU.
- Implementing such a scheme can be **complex**.

Multi Queue Multiprocessor Scheduling (MQMS)

- **MQMS provides each CPU with its own private scheduling queue.**
 - When a job enters the system, it is placed on exactly one scheduling queue.
- **Design Benefits**
 - **Independent Queue Operation:** Each CPU manages and schedules tasks exclusively from its own local queue.
 - **No Synchronization Needed:** Because the queues are independent, CPUs do not need to coordinate access to their respective queues using locks.
 - **Reduced Overhead & Improved Performance:** The removal of lock acquisition and release overhead significantly improves system scalability and overall performance.

MQMS Example


- With Round Robin, the system might produce a schedule that looks like this:




MQMS achieves better performance by reducing inter-core communication (**scalability**) and maximizing local cache utilization (**cache affinity**).

Load Imbalance Issue of MQMS

■ After job C in Q0 finishes:




■ After job A in Q0 finishes:




How to deal with load imbalance?

- The answer is to move jobs (Migration).

- Example:



The OS moves one of B or D to CPU 0



or




How to deal with load imbalance? (Cont.)

- A more tricky case:



- A possible migration pattern:

- Keep switching jobs



CPU Scheduling: Summary

■ Goals

- Efficiency: Maximize CPU utilization
- Fairness: Equal share among processes
- Responsiveness: Minimize response time

■ MLFQ Highlights

- Start at high priority, demote on long use
- Promote periodically to avoid starvation
- Mixes fairness (RR) + efficiency (SJF)

■ Multiprocessor Scheduling

- Cache affinity: keep process on same CPU
- Single Queue (SQMS): simple but low scalability
- Multi-Queue (MQMS): per-CPU queues, use work stealing for load balance

Type	Algorithm
Non-Preemptive	FCFS – simple, causes convoy effect
	SJF – shortest job first, best avg. turnaround
Preemptive	STCF – run job with least time left
	RR – time slicing for fairness
	MLFQ – adaptive priority, learns from behavior